



# ICAPS 2006

The English Lake District, Cumbria, UK

## Tutorial on Planning Graph Based Reachability Heuristics

**Daniel Bryce**  
**Subbarao Kambhampati**  
*Arizona State University*







University of  
HUDDERSFIELD



Carnegie Mellon



RIACS

Honeywell

@ROAD<sup>®</sup>  
MOBILE RESOURCE MANAGEMENT



# icaps

# ICAPS 2006

The English Lake District, Cumbria, UK

## Tutorial on Planning Graph Based Reachability Heuristics

**Daniel Bryce**  
**Subbarao Kambhampati**  
*Arizona State University*





## **ICAPS 2006**

### **Tutorial on Planning Graph Based Reachability Heuristics**

---

#### **Table of contents**

<b>Preface</b>	3
 <b>Presentation</b>	 5
<i>Subbarao Kambhampati, Daniel Bryce</i>	

---

<http://icaps06.icaps-conference.org/>





## ICAPS 2006 Tutorial on Planning Graph Based Reachability Heuristics

---

### Preface

The primary revolution in automated planning in the last decade has been the very impressive scale-up in planner performance. A large part of the credit for this can be attributed squarely to the invention and deployment of powerful reachability heuristics. Most, if not all, modern reachability heuristics are based on a remarkably extensible datastructure called the planning graph—which made its debut as a bit player in the success of Graphplan, but quickly grew in prominence to occupy the center-stage.

In this tutorial, we will start with a discussion of the foundations of reachability analysis with planning graphs. We will then discuss the many ways of applying this analysis to develop scalable planners. Starting with classical planning, we will discuss heuristics for cost-based planning, over-subscription planning, planning with resources, temporal planning, non-deterministic planning as well as stochastic planning.

#### Instructors

- Subbarao Kambhampati  
*rao@asu.edu*  
*http://rakaposhi.eas.asu.edu*  
*Ph: 602-965-0113*  
*Fax: 602-965 2751*  
*Arizona State University*  
*Tempe, AZ 85287-5406*
  
- Daniel Bryce  
*dan.bryce@asu.edu*  
*http://verde.eas.asu.edu*  
*Ph: 602-965-2735*  
*Arizona State University*  
*Tempe, AZ 85287-5406*



# Planning Graph Based Reachability Heuristics

Daniel Bryce & Subbarao Kambhampati

ICAPS'06 Tutorial 6

June 7, 2006



<http://rakaposhi.eas.asu.edu/pg-tutorial/>



[dan.bryce@asu.edu](mailto:dan.bryce@asu.edu)   <http://verde.eas.asu.edu>  
[rao@asu.edu](mailto:rao@asu.edu),        <http://rakaposhi.eas.asu.edu>


## Motivation

- Ways to improve Planner Scalability
  - Problem Formulation
  - Search Space
  - Reachability Heuristics
    - Domain (Formulation) Independent
    - Work for many search spaces
    - Flexible – work with most domain features
    - Overall compliment other scalability techniques
    - Effective!!

# Scalability of Planning

- Before, planning algorithms could synthesize about 6 – 10 action plans in minutes
- Significant scale-up in the last 6-7 years
  - Now, we can synthesize 100 action plans in seconds.

Problem is Search Control!!!



The primary revolution in planning in the recent years has been domain-independent heuristics to scale up plan synthesis

## Topics

- Classical Planning } Rao
- Cost Based Planning } Dan
- Partial Satisfaction Planning } Dan
- Resources (Continuous Quantities) } Rao
- Temporal Planning } Rao
- Non-Deterministic/Probabilistic Planning } Dan
- Hybrid Models } Dan

# Classical Planning

June 7th, 2006


ICAPS'06 Tutorial T6

5

## Rover Domain

```
(define (domain rovers_classical)
  (:requirements :strips :typing)
  (:types waypoint data)
  (:predicates
    (at ?x - waypoint)
    (avail ?d - data ?x - waypoint)
    (comm ?d - data)
    (have ?d - data))
  (:action drive
    :parameters (?x ?y - waypoint)
    :precondition (at ?x)
    :effect (and (at ?y) (not (at ?x))))
  (:action commun
    :parameters (?d - data)
    :precondition (have ?d)
    :effect (comm ?d))
  (:action sample
    :parameters (?d - data ?x - waypoint)
    :precondition (and (at ?x) (avail ?d ?x))
    :effect (have ?d)))
)
```

```
(define (problem rovers_classical1)
  (:domain rovers_classical)
  (:objects
    soil image rock - data
    alpha beta gamma - waypoint)
  (:init (at alpha)
    (avail soil alpha)
    (avail rock beta)
    (avail image gamma))
  (:goal (and (comm soil)
    (comm image)
    (comm rock)))
)
```



June 7th, 2006

ICAPS'06 Tutorial T6

# Classical Planning


- Relaxed Reachability Analysis
- Types of Heuristics
  - Level-based
  - Relaxed Plans
- Mutexes
- Heuristic Search
  - Progression
  - Regression
  - Plan Space
- Exploiting Heuristics

June 7th, 2006

ICAPS'06 Tutorial T6

7

## Planning Graph and Search Tree



- Envelope of Progression Tree (Relaxed Progression)
  - Proposition lists: Union of states at k<sup>th</sup> level
- Lowerbound reachability information

ICAPS'06 Tutorial T6

8

# Level Based Heuristics


- The distance of a proposition is the index of the first proposition layer in which it appears
  - Proposition distance changes when we propagate cost functions – described later
- What is the distance of a Set of propositions??
  - Set-Level: Index of first proposition layer where all goal propositions appear
    - Admissible
    - Gets better with mutexes, otherwise same as max
  - Max: Maximum distance proposition
  - Sum: Summation of proposition distances

June 7th, 2006

ICAPS'06 Tutorial T6

9

## Example of Level Based Heuristics




June 7th, 2006

ICAPS'06 Tutorial T6

10

# Distance of a Set of Literals




- $\text{lev}(p)$  : index of the first level at which p comes into the planning graph
- $\text{lev}(S)$ : index of the first level where all props in S appear non-mutexed.
  - If there is no such level, then
    - If the graph is grown to level off, then  $\infty$
    - Else  $k+1$  ( $k$  is the current length of the graph)

June 7th, 2006

ICAPS'06 Tutorial T6

11

# How do Level-Based Heuristics Break?



June 7th, 2006

ICAPS'06 Tutorial T6

12

# Relaxed Plan Heuristics


- When Level does not reflect distance well, we can find a relaxed plan.
- A relaxed plan is subgraph of the planning graph, where:
  - Every goal proposition is in the relaxed plan at the level where it first appears
  - Every proposition in the relaxed plan has a supporting action in the relaxed plan
  - Every action in the relaxed plan has its preconditions supported.
- Relaxed Plans are not admissible, but are generally effective.
- Finding the optimal relaxed plan is NP-hard, but finding a greedy one is easy. Later we will see how “greedy” can change.

June 7th, 2006

ICAPS'06 Tutorial T6

13

## Example of Relaxed Plan Heuristic




June 7th, 2006

ICAPS'06 Tutorial T6

14


# Results



June 7th, 2006  
ICAPS'06 Tutorial T6

15

# Results (cont'd)





June 7th, 2006  
ICAPS'06 Tutorial T6

16

# Optimizations in Heuristic Computation

- Taming Space/Time costs
  - Bi-level Planning Graph representation
  - Partial expansion of the PG (stop *before* level-off)
    - It is FINE to cut corners when using PG for heuristics (instead of search)!!
- Branching factor can still be quite high
  - Use actions appearing in the PG (complete)
    - Select actions in  $\text{lev}(S)$  vs Levels-off (incomplete)
    - Consider action appearing in RP (incomplete)

June 7th, 2006 ICAPS'06 Tutorial T6




17

# Adjusting for Negative Interactions

- Until now we assume actions only positively interact, but they often conflict
- Mutexes help us capture some negative interactions
  - Types
    - Actions: Interference/Competing Needs
    - Propositions: Inconsistent Support
  - Binary are the most common and practical
  - $|A| + 2|P|$ -ary will allow us to solve the planning problem with a backtrack-free GraphPlan search
    - An action layer may have  $|A|$  actions and  $2|P|$  noops
  - Serial Planning Graph assumes all non-noop actions are mutex

# Binary Mutexes



## Adjusting the Relaxed Plans


- Start with RP heuristic and adjust it to take subgoal interactions into account
  - Negative interactions in terms of “degree of interaction”
  - Positive interactions in terms of co-achievement links
    - Ignore negative interactions when accounting for positive interactions (and *vice versa*)

PROBLEM	Level	Sum	AdjSum2M
Gripper-25	-	69/0.98	67/1.57
Gripper-30	-	81/1.63	77/2.83
Tower-7	127/1.28	127/0.95	127/1.37
Tower-9	511/47.91	511/16.04	511/48.45
8-Puzzle1	31/6.25	39/0.35	31/0.69
8-Puzzle2	30/0.74	34/0.47	30/0.74
Mystery-6	-	-	16/62.5
Mistery-9	8/0.53	8/0.66	8/0.49
Mprime-3	4/1.87	4/1.88	4/1.67
Mprime-4	8/1.83	8/2.34	10/1.49
Alps-grid1	14/1.07	14/1.12	14/0.88
Alps-grid2	-	-	34/95.98

$$HAdjSum2M(S) = \text{length}(\text{RelaxedPlan}(S)) + \max_{p,q \in S} \delta(p,q)$$

Where  $\delta(p,q) = \text{lev}(\{p,q\}) - \max\{\text{lev}(p), \text{lev}(q)\}$  /\*Degree of –ve Interaction \*/

# Anatomy of a State-space Regression planner



Problem: Given a set of subgoals (regressed state)  
estimate how far they are from the initial state


[AAAI 2000; AIPS 2000; AIJ 2002; JAIR 2003]

June 7th, 2006


ICAPS'06 Tutorial T6

21

## Rover Example in Regression



# AltAlt Performance



Problem sets from IPC 2000

June 7th, 2006

ICAPS'06 Tutorial T6

23

# Plan Space Search

Then it was cruelly  
UnPOPped



The good times  
return with Re(vived)POP

In the beginning it was all POP.


June 7th, 2006

ICAPS'06 Tutorial T6

24

## POP Algorithm

<ol style="list-style-type: none"> <li>1. <u>Plan Selection</u>: Select a plan <math>P</math> from the search queue</li> <li>2. <u>Flaw Selection</u>: Choose a flaw <math>f</math> (open cond or unsafe link)</li> <li>3. <u>Flaw resolution</u>:           <ul style="list-style-type: none"> <li>If <math>f</math> is an open condition, choose an action <math>S</math> that achieves <math>f</math></li> <li>If <math>f</math> is an unsafe link, choose promotion or demotion</li> <li>Update <math>P</math></li> <li>Return NULL if no resolution exist</li> </ul> </li> <li>4. If there is no flaw left, return <math>P</math></li> </ol>
---



### Choice points

- Flaw selection (*open condition? unsafe link? Non-backtrack choice?*)
- Flaw resolution/Plan Selection (*how to select (rank) partial plan?*)

June 7th, 2006

ICAPS'06 Tutorial T6



25



## PG Heuristics for Partial Order Planning



- Distance heuristics to estimate cost of partially ordered plans (and to select flaws)
- If we ignore negative interactions, then the set of open conditions can be seen as a regression state
- Mutexes used to detect indirect conflicts in partial plans
  - A step threatens a link if there is a mutex between the link condition and the steps' effect or precondition
  - Post disjunctive precedences and use propagation to simplify




June 7th, 2006

ICAPS'06 Tutorial T6

$S_k \prec S_i \vee S_j \prec S_k$

26

# Regression and Plan Space



## RePOP's Performance

- RePOP implemented on top of UCPOP
  - Dramatically better than any other partial order planner before it
  - Competitive with Graphplan and AltAlt
  - VHPOP carried the torch at ICP 2002

Problem	UCPOP	RePOP	Graphplan	AltAlt
Gripper-8	-	1.01	66.82	.43
Gripper-10	-	2.72	47min	1.15
<b>Gripper-20</b>	<b>-</b>	<b>81.86</b>	<b>-</b>	<b>15.42</b>
Rocket-a	-	8.36	75.12	1.02
<b>Rocket-b</b>	<b>-</b>	<b>8.17</b>	<b>77.48</b>	<b>1.29</b>
Logistics-a	-	3.16	306.12	1.59
Logistics-b	-	2.31	262.64	1.18
Logistics-c	-	22.54	-	4.52
<b>Logistics-d</b>	<b>-</b>	<b>91.53</b>	<b>-</b>	<b>20.62</b>
Bw-large-a	45.78	(5.23) -	14.67	4.12
Bw-large-b	-	(18.86) -	122.56	14.14
<b>Bw-large-c</b>	<b>-</b>	<b>(137.84) -</b>	<b>-</b>	<b>116.34</b>

Written in Lisp, runs on Linux, 500MHz, 250MB

You see, pop, it *is* possible to Re-use all the old POP work!

[IJCAI, 2001]

# Exploiting Planning Graphs

- Restricting Action Choice
  - Use actions from:
    - Last level before level off (complete)
    - Last level before goals (incomplete)
    - First Level of Relaxed Plan (incomplete) – FF's helpful actions
    - Only action sequences in the relaxed plan (incomplete) – YAHSP
- Reducing State Representation
  - Remove static propositions. A static proposition is only ever true or false in the last proposition layer.

# Classical Planning Conclusions

- Many Heuristics
  - Set-Level, Max, Sum, Relaxed Plans
- Heuristics can be improved by adjustments
  - Mutexes
- Useful for many types of search
  - Progression, Regression, POCL

---

# Cost-Based Planning

June 7th, 2006

ICAPS'06 Tutorial T6

31


---

## Cost-based Planning

---

- Propagating Cost Functions
- Cost-based Heuristics
  - Generalized Level-based heuristics
  - Relaxed Plan heuristics

# Rover Cost Model




June 7th, 2006

ICAPS'06 Tutorial T6


33

# Cost Propagation



# Cost Propagation (cont'd)


1-lookahead



## Terminating Cost Propagation

- Stop when:
  - goals are reached (no-lookahead)
  - costs stop changing ( $\infty$ -lookahead)
  - k levels after goals are reached (k-lookahead)

# Guiding Relaxed Plans with Costs



June 7th, 2006

ICAPS'06 Tutorial T6

37

# Cost-Based Planning Conclusions

- Cost-Functions:
  - Remove false assumption that level is correlated with cost
  - Improve planning with non-uniform cost actions
  - Are cheap to compute (constant overhead)

---

## Partial Satisfaction (Over-Subscription) Planning

June 7th, 2006

ICAPS'06 Tutorial T6

39

---


## Partial Satisfaction Planning

---

- Selecting Goal Sets
  - Estimating goal benefit
- Anytime goal set selection
- Adjusting for negative interactions between goals

# Partial Satisfaction (Oversubscription) Planning


In many real world planning tasks, the agent often has more goals than it has resources to accomplish.



## Example: Rover Mission Planning (MER)

Need automated support for  
Over-subscription/Partial Satisfaction  
Planning

Actions have execution costs, goals have utilities, and the objective is to find the plan that has the highest net benefit.




June 7th, 2006

ICAPS'06 Tutorial T6


41

# Adapting PG heuristics for PSP

- **Challenges:**
  - Need to propagate costs on the planning graph
  - The exact set of goals are not clear
    - Interactions between goals
    - Obvious approach of considering all  $2^n$  goal subsets is *infeasible*



- **Idea:** Select a subset of the top level goals upfront
- **Challenge:** Goal interactions
  - Approach: Estimate the net benefit of each goal in terms of its utility minus the cost of its relaxed plan
  - Bias the relaxed plan extraction to (re)use the actions already chosen for other goals




June 7th, 2006

ICAPS'06 Tutorial T6

42

# Goal Set Selection In Rover Problem



June 7th, 2006


ICAPS'06 Tutorial T6

43

## SAPA<sup>PS</sup> (anytime goal selection)

- A\* Progression search
  - g-value: net-benefit of plan so far
  - h-value: relaxed plan estimate of best goal set
    - Relaxed plan found for all goals
    - Iterative goal removal, until net benefit does not increase
  - Returns plans with increasing g-values.

# Some Empirical Results for AltAlt<sup>ps</sup>



Exact algorithms based on MDPs don't scale at all

[AAAI 2004]

June 7th, 2006


ICAPS'06 Tutorial T6

45

## Adjusting for Negative Interactions (AltWlt)

- Problem:
  - What if the apriori goal set is not achievable because of negative interactions?
  - What if greedy algorithm gets bad local optimum?
- Solution:
  - Do not consider mutex goals
  - Add penalty for goals whose relaxed plan has mutexes.
    - Use interaction factor to adjust cost, similar to adjusted sum heuristic
      - $\max_{g_1, g_2 \in G} \{lev(g_1, g_2) - \max(lev(g_1), lev(g_2))\}$
  - Find Best Goal set for each goal

# The Problem with Plangraphs [Smith, ICAPS 04]



**Assume independence between objectives**

For rover: all estimates from starting location

## Approach


- Construct *orienteering* problem
- Solve it
- Use as search guidance

# Orienteering Problem

TSP variant

- Given:
  - network of cities
  - rewards in various cities
  - finite amount of gas
- Objective:
  - collect as much reward as possible
  - before running out of gas

## Orienteering Graph



# The Big Question:

---

How do we determine which propositions go in the orienteering graph?

Propositions that:  
are changed in achieving one goal  
impact the cost of another goal

## Sensitivity analysis


June 7th, 2006

ICAPS'06 Tutorial T6


51

## Recall: Plan Graph


---



# Sensitivity Analysis




# Sensitivity Analysis



## Basis Set Algorithm

```
For each goal:  
    Construct a relaxed plan  
    For each net effect of relaxed plan:  
        Reset costs in PG  
        Set cost of net effect to 0  
        Set cost of mutex initial conditions to  $\infty$   
        Compute revised cost estimates  
        If significantly different,  
            add net effect to basis set
```

## 25 Rocks



# PSP Conclusions

---

- Goal Set Selection
  - Apriori for Regression Search
  - Anytime for Progression Search
  - Both types of search use greedy goal insertion/removal to optimize net-benefit of relaxed plans
- Orienteering Problem
  - Interactions between goals apparent in OP
  - Use solution to OP as heuristic
  - Planning Graphs help define OP

---

# Planning with Resources

# Planning with Resources

- Propagating Resource Intervals
- Relaxed Plans
  - Handling resource subgoals

June 7th, 2006

ICAPS'06 Tutorial T6

59

## Rover with power Resource

```
(define (domain rovers_resource)
(:requirements :strips :typing)
(:types waypoint data)
(:predicates (comm ?d - data)
(have ?d - data)
(at ?x - waypoint)
(avail ?d - data ?x - waypoint))
(:functions (power)
(effort ?x ?y - waypoint)
(effort ?d - data))

(:action drive
:parameters (?x ?y - waypoint)
:precondition (and (at ?x) (>= (power) (effort ?x ?y)))
:effect (and (at ?y) (not (at ?x))
(decrease (power) (effort ?x ?y)))))


(:action commun
:parameters (?d - data)
:precondition (and (have ?d)(>= (power) 5))
:effect (and (comm ?d) (decrease (power) 5)))

(:action sample
:parameters (?d - data ?x - waypoint)
:precondition (and (at ?x) (avail ?d ?x)
(>= (power) (effort ?d)))
:effect (and (have ?d) (decrease (power) (effort ?d)))))

(:action recharge
:parameters ()
:precondition (and (at ?x) (avail ?d ?x) (<= (power) 75))
:effect (and (have ?d) (increase (power) 25)))
)
```

```
(define (problem rovers_resource1)
(:domain rovers_resource)
(:objects
soil image rock - data
alpha beta gamma - waypoint)
(:init (at alpha)
(avail soil alpha)
(avail rock beta)
(avail image gamma)
(= (effort alpha beta) 10)
(= (effort beta alpha) 5)
(= (effort alpha gamma) 30)
(= (effort gamma alpha) 5)
(= (effort beta gamma) 15)
(= (effort gamma beta) 10)
(= (effort soil) 20)
(= (effort rock) 25)
(= (effort image) 5)
(= (power) 25))
(:goal (and (comm soil)
(comm image)
(comm rock))))
```

Resource Usage,  
Same as costs for  
This example




June 7th, 2006


ICAPS'06 Tutorial T6

60


# Resource Intervals



# Resource Intervals (cont'd)



# Relaxed Plan Extraction with Resources



June 7th, 2006

ICAPS'06 Tutorial T6

63

# Results




Figure 13: Runtime curves on Settlers instances for the planners favoring speed. Time is shown on a logarithmic scale, instance size scales from left to right.

## Results (cont'd)




Figure 8: Runtime curves on *Depots* instances for the planners favoring speed. Time is shown on a logarithmic scale, instance size scales from left to right.

## Planning With Resources Conclusion

- Resource Intervals allow us to be optimistic about reachable values
  - Upper/Lower bounds can get large
- Relaxed Plans may require multiple supporters for subgoals
- Negative Interactions are much harder to capture

---

# Temporal Planning

June 7th, 2006

ICAPS'06 Tutorial T6

67


---

# Temporal Planning

---

- Temporal Planning Graph
  - From Levels to Time Points
  - Delayed Effects
- Estimating Makespan
- Relaxed Plan Extraction

# Rover with Durative Actions




June 7th, 2006

ICAPS'06 Tutorial T6

69

# SAPA




[ECP 2001; AIPS 2002; ICAPS 2003; JAIR 2003]

## Search through time-stamped states

### Goal Satisfaction:

$S = (P, M, \Pi, Q, t) \Rightarrow G$  if  $\forall \langle p_i, t_i \rangle \in G$  either:

- $\exists \langle p_i, t_j \rangle \in P, t_j < t_i$  and no event in  $Q$  deletes  $p_i$ .
- $\exists e \in Q$  that adds  $p_i$  at time  $t_e < t_i$ .




### Action Application:

Action  $A$  is applicable in  $S$  if:

- All instantaneous preconditions of  $A$  are satisfied by  $P$  and  $M$ .
- $A$ 's effects do not interfere with  $\Pi$  and  $Q$ .
- No event in  $Q$  interferes with persistent preconditions of  $A$ .
- $A$  does not lead to concurrent resource change

### When $A$ is applied to $S$ :


- $P$  is updated according to  $A$ 's instantaneous effects.
- Persistent preconditions of  $A$  are put in  $\Pi$
- Delayed effects of  $A$  are put in  $Q$ .




### Search:

Pick a state  $S$  from the queue.  
 If  $S$  satisfies the goals, end  
 Else non-deterministically do one of  
 --Advance the clock  
 (by executing the earliest event in  $Q_s$ )  
 --Apply one of the applicable actions to  $S$

## Temporal Planning



# SAPA at IPC-2002



June 7th, 2006

ICAPS'06 Tutorial T6

[JAIR 2003] 73

## Temporal Planning Conclusion

- Levels become Time Points
- Makespan and plan length/cost are different objectives
- Set-Level heuristic measures makespan
- Relaxed Plans measure makespan and plan cost

---

## Non-Deterministic Planning

June 7th, 2006

ICAPS'06 Tutorial T6

75

---

## Non-Deterministic Planning


---

- Belief State Distance
- Multiple Planning Graphs
- Labelled Uncertainty Graph
- Implicit Belief states and the CFF heuristic

# Conformant Rover Problem

```
(define (domain rovers_conformant)
  (:requirements :strips :typing)
  (:types waypoint data)
  (:predicates
    (at ?x - waypoint)
    (avail ?d - data ?x - waypoint)
    (comm ?d - data)
    (have ?d - data))
  (:action drive
    :parameters (?x ?y - waypoint)
    :precondition (at ?x)
    :effect (and (at ?y) (not (at ?x))))
  (:action commun
    :parameters (?d - data)
    :precondition (have ?d)
    :effect (comm ?d))
  (:action sample
    :parameters (?d - data ?x - waypoint)
    :precondition (at ?x)
    :effect (when (avail ?d ?x) (have ?d)))
  )
```

```
(define (problem rovers_conformant1)
  (:domain rovers)
  (:objects
    soil image rock - data
    alpha beta gamma - waypoint)
  (:init (at alpha)
    (oneof (avail soil alpha)
      (avail soil beta)
      (avail soil gamma)))
  (:goal (comm soil))
  )
```




June 7th, 2006

ICAPS'06 Tutorial T6

77

# Search in Belief State Space




June 7th, 2006


ICAPS'06 Tutorial T6

78

# Belief State Distance



# Belief State Distance



# State Distance Aggregations

[Bryce et.al, 2005]





Figure 16: Ratio of heuristic estimates for distance between  $BS_I$  and  $BS_G$  to optimal plan length. Rv = Rovers, L = Logistics, B = BT, BC = BTC, C = Cube Center, R = Ring.

June

81


# Multiple Planning Graphs

Extract Relaxed Plans from each  
Step-wise union relaxed plans




# Labelled Uncertainty Graph

Action labels are the Conjunction (intersection) of their Precondition labels



# Labelled Relaxed Plan



# Comparison of Planning Graph Types

[JAIR, 2006]




Figure 15: Total Time (ms) to solve all problems when sampling worlds to use in heuristic computation.

June 7th, 2006

ICAPS'06 Tutorial T6


85

## State Agnostic Planning Graphs (SAG)

- LUG represents multiple explicit planning graphs
- SAG uses LUG to represent a planning graph for every state
- The SAG is built once per search episode and we can use it for relaxed plans for every search node, instead of building a LUG at every node
- Extract relaxed plans from SAG by ignoring planning graph components not labeled by states in our search node.

# SAG

**Build a LUG for all states (union of all belief states)**




[ AAAI, 2005]

June 7th, 2006

ICAPS'06 Tutorial T6

87

# CFF (implicit belief states + PG)




June 7th, 2006


ICAPS'06 Tutorial T6

88

## Belief Space Problems




## Classical Problems



Conformant

More Results at the IPC!!!



## Conditional Planning

- Actions have Observations
- Observations branch the plan because:
  - Plan Cost is reduced by performing less “just in case” actions – each branch performs relevant actions
  - Sometimes actions conflict and observing determines which to execute (e.g., medical treatments)
- We are ignoring negative interactions
  - We are only forced to use observations to remove negative interactions
  - Ignore the observations and use the conformant relaxed plan
    - Suitable because the aggregate search effort over all plan branches is related to the conformant relaxed plan cost

# Non-Deterministic Planning Conclusions

- Measure positive interaction and independence between states co-transitioning to the goal via overlap
  - Labeled planning graphs and CFF SAT encoding efficiently measure conformant plan distance
- Conformant planning heuristics work for conditional planning without modification

June 7th, 2006

ICAPS'06 Tutorial T6

91


# Stochastic Planning

# Stochastic Rover Example

[ICAPS 2006]

```
(define (domain rovers_stochastic)
  (:requirements :strips :typing)
  (:types waypoint data)
  (:predicates
    (at ?x - waypoint)
    (avail ?d - data ?x - waypoint)
    (comm ?d - data)
    (have ?d - data))
  (:action drive
    :parameters (?x ?y - waypoint)
    :precondition (at ?x)
    :effect (and (at ?y) (not (at ?x))))
  (:action commun
    :parameters (?d - data)
    :precondition (have ?d)
    :effect (probabilistic 0.8 (comm ?d)))
  (:action sample
    :parameters (?d - data ?x - waypoint)
    :precondition (at ?x)
    :effect (when (avail ?d ?x)
      (probabilistic 0.9 (have ?d))))
  )
```

```
(define (problem rovers_stochastic1)
  (:domain rovers)
  (:objects
    soil image rock - data
    alpha beta gamma - waypoint)
  (:init (at alpha))
  (:probabilistic 0.4 (avail soil alpha)
    0.5 (avail soil beta)
    0.1 (avail soil gamma))
  (:goal (comm soil) 0.5)
  )
```




June 7th, 2006

ICAPS'06 Tutorial T6

93

# Search in Probabilistic Belief State Space



June 7th, 2006

ICAPS'06 Tutorial T6

94

# Handling Uncertain Actions

[ICAPS 2006]

- Extending LUG to handle uncertain actions requires label extension that captures:
  - State uncertainty (as before)
  - Action outcome uncertainty
    - Problem: Each action at each level may have a different outcome. The number of uncertain events grows over time – meaning the number of joint outcomes of events grows exponentially with time
    - Solution: Not all outcomes are important. Sample some of them – keep number of joint outcomes constant.


# Monte Carlo LUG (McLUG)

[ICAPS 2006]

- Use Sequential Monte Carlo in the Relaxed Planning Space
  - Build several deterministic planning graphs by sampling states and action outcomes
  - Represent set of planning graphs using LUG techniques
    - Labels are sets of particles
    - Sample which Action outcomes get labeled with particles
    - Bias relaxed plan by picking actions labeled with most particles to prefer more probable support

# McLUG for rover example

[ICAPS 2006]




June 7th, 2006

ICAPS'06 Tutorial T6

97

# Logistics Domain Results

[ICAPS 2006]




June 7th, 2006

ICAPS'06 Tutorial T6



98

# Grid Domain Results


[ICAPS 2006]



Again, good scalability and quality!



Need More Particles for broad beliefs



June 7th, 2006

99

# Direct Probability Propagation

- Alternative to label propagation, we can propagate numeric probabilities
  - Problem: Numeric Propagation tends to assume only independence or positive interaction between actions and propositions.
    - With probability, we can vastly under-estimate the probability of reaching propositions
  - Solution: Propagate Correlation – measures pair-wise independence/pos interaction/neg interaction
    - Can be seen as a continuous mutex

# Correlation

---

- $C(x, y) = \Pr(x, y)/(\Pr(x)\Pr(y))$
- If :
  - $C(x, y) = 0$ , then  $x, y$  are mutex
  - $0 < C(x, y) < 1$ , then  $x, y$  interfere
  - $C(x, y) = 1$ , then  $x, y$  are independent
  - $1 < C(x, y) < 1/\Pr(x)$ , then  $x, y$  synergize
  - $C(x, y) = 1/\Pr(x) = 1/\Pr(y)$ , then  $x, y$  are completely correlated

# Probability of a set of Propositions

---

- $\Pr(x_1, x_2, \dots, x_n) = \prod_{i=1..n} \Pr(x_i | x_1 \dots x_{i-1})$  **Chain Rule**
  - $\Pr(x_i | x_1 \dots x_{i-1}) = \frac{\Pr(x_1 \dots x_{i-1} | x_i) \Pr(x_i)}{\Pr(x_1 \dots x_{i-1})}$  **Bayes Rule**

$$\approx \frac{\Pr(x_1 | x_i) \dots \Pr(x_{i-1} | x_i) \Pr(x_i)}{\Pr(x_1) \dots \Pr(x_{i-1})}$$
 **Assume Independence**

$$= \frac{\Pr(x_i | x_1)}{\Pr(x_i)} \dots \frac{\Pr(x_i | x_{i-1})}{\Pr(x_i)} \Pr(x_i)$$
 **Bayes Rule**

$$= \Pr(x_i) C(x_i, x_1) \dots C(x_i, x_{i-1})$$
 **Correlation**

$$= \Pr(x_i) \prod_{j=1..i-1} C(x_i, x_j)$$
- $\Pr(x_1, x_2, \dots, x_n) = \prod_{i=1..n} \Pr(x_i) \prod_{j=1..i-1} C(x_i, x_j)$

# Probability Propagation

- The probability of an Action being enabled is the probability of its preconditions (a set of propositions).
- The probability of an effect is the product of the action probability and outcome probability
- A single (or pair of) proposition(s) has probability equal to the probability it is given by the best set of supporters.
- The probability that a set of supporters gives a proposition is the sum over the probability of all possible executions of the actions.

June 7th, 2006

ICAPS'06 Tutorial T6

103

## Results




Figure 2: Run times (s), Plan lengths, and Expanded Nodes vs. probability threshold for sandcastle-67




Figure 3: Run times (s), Plan lengths, and Expanded Nodes vs. probability threshold for slippery gripper

June 7th, 2006

ICAPS'06 Tutorial T6

104

# Stochastic Planning Conclusions

- Number of joint action outcomes too large
  - Sampling outcomes to represent in labels is much faster than exact representation
- SMC gives us a good way to use multiple planning graph for heuristics, and the McLUG helps keep the representation small
- Numeric Propagation of probability can better capture interactions with correlation
  - Can extend to cost and resource propagation

June 7th, 2006

ICAPS'06 Tutorial T6

105

# Hybrid Planning Models

# Hybrid Models




- Metric-Temporal w/ Resources (SAPA)
- Temporal Planning Graph w/ Uncertainty (Prottle)
- PSP w/ Resources (SAPA<sup>MPS</sup>)
- Cost-based Conditional Planning (CLUG)

June 7th, 2006

ICAPS'06 Tutorial T6

107

# Propagating Temporal Cost Functions



```

Shuttle(Tempe,Phx):
Cost: $20; Time: 1.0 hour
Helicopter(Tempe,Phx):
Cost: $100; Time: 0.5 hour
Car(Tempe,LA):
Cost: $100; Time: 10 hour
Airplane(Phx,LA):
Cost: $200; Time: 1.0 hour

```

## Heuristics based on cost functions


### Direct

- If we want to minimize makespan:
  - $h = t_0$
- If we want to minimize cost
  - $h = \text{CostAggregate}(G, t_\infty)$
- If we want to minimize a function  $f(\text{time}, \text{cost})$  of cost and makespan
  - $h = \min f(t, \text{Cost}(G, t)) \text{ s.t. } t_0 \leq t \leq t_\infty$ 
    - E.g.  $f(\text{time}, \text{cost}) = 100 \cdot \text{makespan} + \text{Cost}$  then  $h = 100x2 + 220$  at  $t_0 \leq t = 2 \leq t_\infty$

### Using Relaxed Plan

- Extract a relaxed plan using  $h$  as the bias
- If the objective function is  $f(\text{time}, \text{cost})$ , then action **A** (to be added to RP) is selected such that:
 
$$f(t(\text{RP}+A), C(\text{RP}+A)) + f(t(G_{\text{new}}), C(G_{\text{new}}))$$
 is minimal

$$G_{\text{new}} = (G \cup \text{Precond}(A)) \setminus \text{Effects}$$



June 7th, 2006

ICAPS'06 Tutorial T6

109

## Phased Relaxation

The relaxed plan can be adjusted to take into account constraints that were originally ignored

### Adjusting for Mutexes:

Adjust the make-span estimate of the relaxed plan by marking actions that are mutex (and thus cannot be executed concurrently)

### Adjusting for Resource Interactions:

Estimate the number of additional resource-producing actions needed to make-up for any resource short-fall in the relaxed plan



$$C = C + \sum_R \lceil (Con(R) - (Init(R) + Pro(R))) / \Delta_R \rceil * C(A_R)$$

June 7th, 2006

ICAPS'06 Tutorial T6

110

# Handling Cost/Makespan Tradeoffs



Results over 20 randomly generated temporal logistics problems involve moving 4 packages between different locations in 3 cities:

$$O = f(\text{time}, \text{cost}) = \alpha \cdot \text{Makespan} + (1 - \alpha) \cdot \text{TotalCost}$$


## Prottle

- SAPA-style (time-stamped states and event queues) search for fully-observable conditional plans using L-RTDP
- Optimize probability of goal satisfaction within a given, finite makespan
- Heuristic estimates probability of goal satisfaction in the plan suffix

# Prottle planning graph

Time Stamped Propositions

Probabilistic Outcomes at Different Times



June 7th, 2006

ICAPS'06 Tutorial T6

113

# Probability Back-Propagation


Heuristic is product of Relevant Node costs

Vector of probability Costs for each top-level Goal:  
 $1 - \Pr(\text{Satisfy } g_i)$

Outcome cost is max of affected proposition costs

Action Cost is expectation of outcome costs

Proposition cost is Product of supported Action costs



June 7th, 2006

ICAPS'06 Tutorial T6

114

# Prottle Results


problem	horizon	$\epsilon$	time1	time2	cost1	cost2	states1	states2
AI	100	0.3	-	103	-	0.344	-	346,100
AI	120	0.6	-	404	-	0.222	-	1,319,229
MS	15	0.0	-	272	-	0.027	-	496,096
MS	15	0.1	-	171	-	0.114	-	309,826
MS	15	0.2	2,431	21	0.119	0.278	13,627,753	6,759
MS	15	0.3	367	235	0.278	0.278	1,950,134	434,772
MZ	10	0.0	195	10	0.178	0.178	1,374,541	13,037
MZ	10	0.1	185	2	0.193	0.178	1,246,159	2,419
MZ	10	0.2	64	1	0.197	0.193	436,876	669
MZ	10	0.3	62	2	0.202	0.193	414,414	1,812
TP	20	0.0	442	< 1	0.798	0.798	3,565,698	3,676
TP	20	0.1	456	< 1	0.798	0.798	3,628,300	2,055
TP	20	0.2	465	< 1	0.798	0.798	3,672,348	2,068
TP	20	0.3	464	< 1	1.000	0.798	3,626,404	1,256

# PSP w/ Resources

- Utility and Cost based on the values of resources
- Challenges:
  - Need to propagate cost for resource intervals
  - Need to support resource goals at different levels


# Resource Cost Propagation

- Propagate reachable values with cost




## Cost Propagation on Variable Bounds

- Bound cost dependent upon
  - action cost
  - previous bound cost
    - current bound cost adds to the next
  - Cost of all bounds in expressions



## Results – *Modified Rovers* (numeric soil)



Average improvement: 3.06


June 7th, 2006

ICAPS'06 Tutorial T6

119

## Anytime A\* Search Behavior

Rovers Utility, Problem 11




June 7th, 2006

ICAPS'06 Tutorial T6

120

## Results – *Modified Logistics* (#of packages)



June 7th, 2006


ICAPS'06 Tutorial T6

121

## Cost-Based Conditional Planning





- Actions may reduce uncertainty, but cost a lot
  - Do we want more “just in case” actions that are cheap, or less that are more expensive
- Propagate Costs on the LUG (CLUG)
  - Problem: LUG represents multiple explicit planning graphs and the costs can be different in each planning graph.
    - A single cost for every explicit planning assumes full positive interaction
    - Multiple costs, one for each planning graph is too costly
  - Solution: Propagate cost for partitions of the explicit planning graphs

# Cost Propagation on the LUG (CLUG)



# The Medical Specialist [Bryce & Kambhampati, 2005]

# Using Propagated Costs Improves Plan Quality, Without much additional time



# Overall Conclusions

- Relaxed Reachability Analysis
  - Concentrate strongly on positive interactions and independence by ignoring negative interaction
  - Estimates improve with more negative interactions
- Heuristics can estimate and aggregate costs of goals or find relaxed plans
- Propagate numeric information to adjust estimates
  - Cost, Resources, Probability, Time
- Solving hybrid problems is hard
  - Extra Approximations
  - Phased Relaxation
  - Adjustments/Penalties

June 7th, 2006

ICAPS'06 Tutorial T6

125

# Why do we love PG Heuristics?

- They work!
- They are “forgiving”
  - You don't like doing mutex? okay
  - You don't like growing the graph all the way? okay.
- Allow propagation of many types of information
  - Level, subgoal interaction, time, cost, world support, probability
- Support phased relaxation
  - E.g. Ignore mutexes and resources and bring them back later...
- Graph structure supports other synergistic uses
  - e.g. action selection
- Versatility...

# Versatility of PG Heuristics

- PG Variations
  - Serial
  - Parallel
  - Temporal
  - Labelled

- Propagation Methods
  - Level
  - Mutex
  - Cost
  - Label

- Planning Problems
  - Classical
  - Resource/Temporal
  - Conformant

- Planners
  - Regression
  - Progression
  - Partial Order
  - Graphplan-style

## References

1. J. Benton, M.B. Do, and S. Kambhampati. "Oversubscription planning with metric goals." In Proceedings of IJCAI'05, 2005.
2. A. Blum and M. Furst. "Fast planning through planning graph analysis." In Proceedings of IJCAI'95, 1995.
3. B. Bonet and H. Geffner. "Planning as heuristic search: New results." In Proceedings of ECP'99, 1999.
4. R. Brafman and J. Hoffmann. "Contingent planning via heuristic forward search with implicit belief states." In Proceedings of ICAPS'05, 2005.
5. D. Bryce and S. Kambhampati. "Cost sensitive reachability heuristics for handling state uncertainty." In Proceedings of UAI'05, 2005.
6. D. Bryce and D.E. Smith. "Using correlation to compute better probability estimates in plan graphs." In ICAPS'06 Workshop on Planning Under Uncertainty and Execution Control for Autonomous Systems, 2006.
7. D. Bryce, S. Kambhampati, and D.E. Smith. "Sequential monte carlo in probabilistic planning reachability heuristics." In Proceedings of ICAPS'06, 2006.
8. D. Bryce, S. Kambhampati, and D.E. Smith. "Planning graph heuristics for belief space search." Journal of Artificial Intelligence Research, to appear.
9. M. Cayrol, P. Rgnier, and V. Vidal. "New results about lcp, a least committed graphplan." In Proceedings of AIPS'00, 2000.
10. W. Cushing and D. Bryce. "State agnostic planning graphs." In Proceedings of AAAI'05, 2005.
11. M.B. Do and S. Kambhampati. "Sapa: A scalable multi-objective metric temporal planner." Journal of Artificial Intelligence Research, 20:155-194, 2003.
12. M. B. Do and S. Kambhampati. "Partial satisfaction (over-subscription) planning as heuristic search." In Proceedings of KBCS'04, 2004.
13. A. Gerevini, A. Saetti, and I. Serina. "Planning through stochastic local search and temporal action graphs in lpg." Journal of Artificial Intelligence Research, 20:239-290, 2003.
14. A. Gerevini, A. Saetti, and I. Serina. "Integrating planning and temporal reasoning for domains with durations and time windows." In Proceedings of IJCAI'05, 2005.
15. M. Ghallab and H. Laruelle. "Representation and control in Ixtet." In Proceedings of AIPS'94, 1994.
16. M. Ghallab, D. Nau, and P. Traverso. "Automated Planning: Theory and Practice." Morgan Kaufmann, 2004.
17. P. Haslum and H. Geffner. Admissible heuristics for optimal planning. In Proceedings of AIPS'00, 2000.
18. M. Helmert. "A planning heuristic based on causal graph analysis." In Proceedings of ICAPS'04, 2004.
19. J. Hoffmann and R. Brafman. "Conformant planning via heuristic forward search: A new approach." In Proceedings of ICAPS'04, 2004.
20. J. Hoffmann and H. Geffner. Branching matters: Alternative branching in graphplan. In Proceedings of ICAPS'03, 2003.
21. J. Hoffmann and B. Nebel. "The FF planning system: Fast plan generation through heuristic search." Journal of Artificial Intelligence Research, 14:253-302, 2001.

22. S. Kambhampati and R. Sanchez. "Distance based goal ordering heuristics for graphplan." In Proceedings of AIPS'00, 2000.
23. S. Kambhampati, E. Lambrecht, and E. Parker. "Understanding and extending graphplan." In Proceedings of ECP'97, 1997.
24. I. Little, D. Aberdeen, and S. Theibaux. "Prottle: A probabilistic temporal planner." In Proceedings of AAAI'05, 2005.
25. D. McDermott. "A heuristic estimator for means-ends analysis in planning." In Proceedings of AIPS'96, 1996.
26. D. McDermott. "Pddl-the planning domain definition language." Technical report, Available at: [www.cs.yale.edu/homes/dvm](http://www.cs.yale.edu/homes/dvm), 1998.
27. D. McDermott. "Using regression-match graphs to control search in planning." Artificial Intelligence, 109(1-2):111-159, 1999.
28. X. Nguyen, S. Kambhampati, and R.S. Nigenda. "Planning graph as the basis for deriving heuristics for plan synthesis by state space and CSP search." Artificial Intelligence, 135(1-2):73-123, 2002.
29. I. Refanidis and I. Vlahavas. "The GRT planning system: Backward heuristic construction in forward statespace planning." Journal of Artificial Intelligence Research, 15:115161, 2001.
30. R. Sanchez and S. Kambhampati. "Planning graph heuristics for selecting objectives in over-subscription planning problems." In Proceedings of ICAPS'05, 2005.
31. D. E. Smith. "Choosing objectives in over-subscription planning." In Proceedings of ICAPS'04, 2004.
32. M. van den Briel, R. Sanchez, M.B. Do, and S. Kambhampati. "Effective approaches for partial satisfaction (over-subscription) planning." In Proceedings of AAAI'04, 2004.
33. V. Vidal. "A lookahead strategy for heuristic search planning." In Proceedings of ICAPS'04, 2004.
34. S. Yoon, A. Fern, and R. Givan. "Learning heuristic functions from relaxed plans." In Proceedings of ICAPS'06, 2006.
35. Terry Zimmerman and Subbarao Kambhampati. "Using memory to transform search on the planning graph." JAIR, 23, 2005.

